

Random Forest

Part 3: with R package randomForest

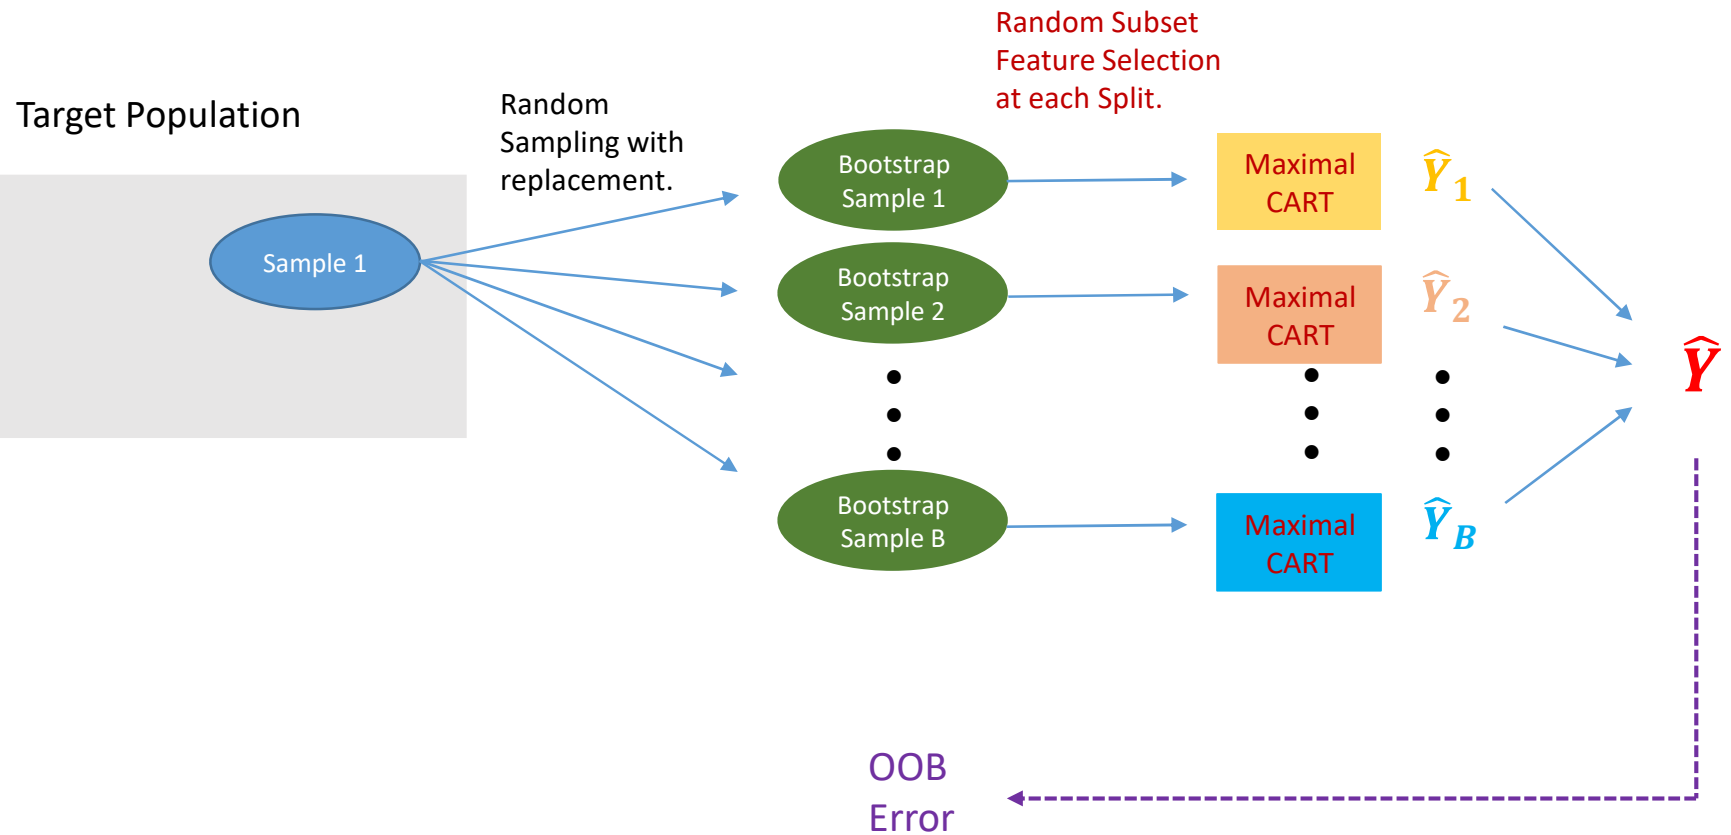
Random Forest with R

- What are the default settings for B and RSF size in R?
- Where do we see OOB error?
- How do we check if the error had converged?
*“...in practice we use **a value of B sufficiently large for the error rate to have settled down.**” -- ISLR*
- What are the useful charts/tables and diagnostics in R?
- What to watch out for if we use Python implementation instead of R?



Random Forest = Bagging +
Random Subset Feature

Random Forest Process



R settings

B = 500 [new default in R]

Categorical Y

- RSF size = $\text{floor}(\text{sqrt}(M))$

Continuous Y

- RSF size = $\text{floor}(M/3)$

Notes:

- M is the number of X variables.
- The floor() function is the greatest integer smaller than the value.
- *Examples: $\text{floor}(3.1) = 3$, $\text{floor}(3.9) = 3$.*
- For categorical Y, Breiman(2001) proposed $\text{int}(\log_2(M) + 1)$. Actually, not much difference if M is a small number.
- For continuous Y, correlation betw CARTs is less sensitive to RSF size.
“Correlation increases quite slowly as the number of features used increases.” --- Breiman(2001)
- *Hence, we allow bigger RSF size for continuous Y as risk of highly correlated CARTS is lower.*
- *Examples:*
 - If $M = 10$, $\text{floor}(\text{sqrt}(10)) = \text{floor}(3.16...) = 3$; $\text{floor}(10/3) = \text{floor}(3.33...) = 3$.
 - If $M = 100$, $\text{floor}(\text{sqrt}(100)) = \text{floor}(10) = 10$; $\text{floor}(100/3) = \text{floor}(33.33...) = 33$!

Comparisons of Different RSF Size Formula Outputs

| M | For Categorical Y | | For Continuous Y |
|------|-----------------------------|------------------------------|-------------------|
| | $\text{INT}(\log_2(M) + 1)$ | $\text{INT}(\text{sqrt}(M))$ | $\text{INT}(M/3)$ |
| 5 | 3 | 2 | 1 |
| 10 | 4 | 3 | 3 |
| 15 | 4 | 3 | 5 |
| 20 | 5 | 4 | 6 |
| 25 | 5 | 5 | 8 |
| 30 | 5 | 5 | 10 |
| 50 | 6 | 7 | 16 |
| 100 | 7 | 10 | 33 |
| 500 | 9 | 22 | 166 |
| 1000 | 10 | 31 | 333 |

Dataset: Heart.csv

Source: <https://archive.ics.uci.edu/ml/datasets/Heart+Disease>

- 303 cases
- 13 Xs and 1 categorical Y (AHD)
- $M = 13$ and RSF size = $\text{floor}(\sqrt{13}) = \text{floor}(3.605...) = 3$
- 6 missing values

| | A | B | C | D | E | F | G | H | I | J | K | L | M | N |
|---|-----|-----|--------------|--------|------|-----|---------|-------|-------|---------|-------|----|------------|-----|
| 1 | Age | Sex | ChestPain | RestBP | Chol | Fbs | RestECG | MaxHR | ExAng | Oldpeak | Slope | Ca | Thal | AHD |
| 2 | 63 | 1 | typical | 145 | 233 | 1 | 2 | 150 | 0 | 2.3 | 3 | 0 | fixed | No |
| 3 | 67 | 1 | asymptomatic | 160 | 286 | 0 | 2 | 108 | 1 | 1.5 | 2 | 3 | normal | Yes |
| 4 | 67 | 1 | asymptomatic | 120 | 229 | 0 | 2 | 129 | 1 | 2.6 | 2 | 2 | reversable | Yes |
| 5 | 37 | 1 | nonanginal | 130 | 250 | 0 | 0 | 187 | 0 | 3.5 | 3 | 0 | normal | No |
| 6 | 41 | 0 | nontypical | 130 | 204 | 0 | 2 | 172 | 0 | 1.4 | 1 | 0 | normal | No |
| 7 | 56 | 1 | nontypical | 120 | 226 | 0 | 0 | 178 | 0 | 0.8 | 1 | 0 | normal | No |

Data Dictionary: Heart Data Dictionary.txt

Rscript: RF Heart1.R

```
5 library(randomForest)
6
7 setwd("C:/NC/Datasets/ML")
8
9 heart.df <- read.csv("Heart.csv", stringsAsFactors = T)
10
11 sum(is.na(heart.df))
12 ## 6 missing values. Need to explicitly handle these in randomForest().
13 ## Options: na.action = na.omit or na.action = na.roughfix
14
15 set.seed(1) # for Bootstrap sampling & RSF selection.
16
17 m.RF.1 <- randomForest(AHD ~ . , data = heart.df,
18                       na.action = na.omit,
19                       importance = T)
20
21 m.RF.1 ## shows defaults are B = 500, RSF size = int(sqrt(m)) = 3
22
23 var.impt <- importance(m.RF.1)
24
25 varImpPlot(m.RF.1, type = 1)
```


?randomForest() at console or open randomForest.PDF for documentation

Usage

```
## S3 method for class 'formula'
randomForest(formula, data=NULL, ..., subset, na.action=na.fail)
## Default S3 method:
randomForest(x, y=NULL, xtest=NULL, ytest=NULL, ntree=500,
  mtry=if (!is.null(y) && !is.factor(y))
    max(floor(ncol(x)/3), 1) else floor(sqrt(ncol(x))),
  replace=TRUE, classwt=NULL, cutoff, strata,
  sampsize = if (replace) nrow(x) else ceiling(.632*nrow(x)),
  nodesize = if (!is.null(y) && !is.factor(y)) 5 else 1,
  maxnodes = NULL,
  importance=FALSE, localImp=FALSE, nPerm=1,
  proximity, oob.prox=proximity,
  norm.votes=TRUE, do.trace=FALSE,
  keep.forest=!is.null(y) && is.null(xtest), corr.bias=FALSE,
  keep.inbag=FALSE, ...)
```

To overwrite if there are missing values. **na.omit** or **na.roughfix**

RSF size

Min cases in
terminal node

Overwrite to
T to get
Variable
Importance.

B

Results of Random Forest on Heart data


```
call:
 randomForest(formula = AHD ~ ., data = heart.df, importance = T,      na.action = na.omit)
      Type of random forest: classification
      Number of trees: 500
No. of variables tried at each split: 3

      OOB estimate of  error rate: 18.18%
Confusion matrix:
      No Yes class.error
No  136  24   0.1500000
Yes   30 107   0.2189781
```

- $B = 500$, RSF size = 3; OOB overall error = $(24 + 30)/297 = 18.18\%$.
- Q: How are the confusion matrix results determined?
 - Ans: From OOB data and majority rule.
- Q: Why did the confusion matrix contain only 297 cases when the dataset has 303 cases?
 - Ans: 6 cases has missing values and were omitted as `na.action= na.omit`

View how OOB error rates change with different number of trees with `plot(m.RF.1)`

OOB Error Rates of Random Forest on Heart data up till 500 Trees



Errors stabilised after approx. 250 trees.

i.e. errors would be approx. the same if `ntree` = any number bigger than 250.

If errors are still exhibiting decreasing trend at 500 trees, set `ntree` to be a bigger number > 500 and view the plot to check stability of errors.

View RF prediction for each case via m.RF.1\$predicted

```
> m.RF.1$predicted
```

| | | | | | | | | | | | | | | | | | | | | | | |
|----|-----|-----|----|----|----|-----|----|-----|-----|----|----|-----|----|----|----|----|----|----|----|----|----|----|
| 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 | 20 | 21 | 22 | 23 |
| No | Yes | Yes | No | No | No | Yes | No | Yes | Yes | No | No | Yes | No | No | No | No | No | No | No | No | No | No |
| 24 | 25 | 26 | 27 | 28 | 29 | 30 | 31 | 32 | 33 | 34 | 35 | 36 | 37 | 38 | 39 | 40 | 41 | 42 | 43 | 44 | 45 | 46 |

To check how many times each case is OOB among the 500 trees, `m.RF.1$oob.times`

$$P(\text{case } i \text{ is OOB}) = (1 - 1/n)^n = (1 - 1/297)^{297} \approx 0.367$$

```
> m.RF.1$oob.times
 [1] 177 185 193 174 174 180 196 169 172 196 192 176 190 189
[15] 192 205 172 194 182 193 183 176 191 185 204 187 201 190
[29] 203 194 172 203 185 166 201 171 187 185 196 196 177 173
[43] 181 184 208 191 166 192 178 184 184 175 190 169 182 193
[57] 209 194 174 176 160 187 201 172 198 177 181 171 174 181
[71] 197 182 182 173 189 159 195 185 187 175 196 197 176 167
[85] 178 196 199 177 172 201 195 184 187 183 178 184 205 191
[99] 168 190 192 181 185 180 168 188 174 182 202 184 189 186
[113] 183 193 177 204 193 176 190 194 179 175 190 215 165 194
[127] 183 192 167 164 176 186 158 204 180 166 196 180 172 175
[141] 180 180 206 196 178 178 193 176 186 184 171 178 179 186
[155] 195 194 183 161 184 167 165 189 186 181 180 169 186 178
[169] 191 180 178 199 191 174 184 205 168 171 194 184 189 196
[183] 176 191 184 190 203 183 183 193 181 179 182 194 164 171
[197] 189 171 183 172 175 167 198 179 176 185 191 211 189 174
[211] 192 164 187 194 171 191 186 183 177 191 183 198 176 177
[225] 175 163 179 191 185 189 190 193 175 179 202 177 189 182
[239] 177 172 186 181 198 196 188 185 159 201 199 171 181 188
[253] 195 194 187 188 192 205 190 175 172 182 181 173 180 192
[267] 178 172 189 173 181 183 193 187 178 177 181 183 181 173
[281] 174 190 197 182 200 175 192 181 201 206 152 187 179 176
[295] 169 181 183
```

Checking:

Case 1: $P(\text{OOB}) = 177/500 = 0.354$

Case 2: $P(\text{OOB}) = 185/500 = 0.37$

Case 3: $P(\text{OOB}) = 193/500 = 0.386$

...

Interpretation:

case 1 is not inside 177 of the 500 trees in the forest, ...

To check whether each case is inbag or OOB in which tree, override the parameter `keep.inbag = T` in `randomForest()`

```
m.RF.1 <- randomForest(AHD ~ . , data=heart.df,  
                        na.action=na.omit,  
                        importance=T,  
                        keep.inbag = T)
```

```
> inbag <- m.RF.1$inbag  
> view(inbag)  
> nrow(inbag)  
[1] 297
```

| | V1 | V2 | V3 | V4 | V5 | V6 | V7 | V8 | V9 | V10 | V11 | V12 | V13 |
|---|----|----|----|----|----|----|----|----|----|-----|-----|-----|-----|
| 1 | 0 | 1 | 2 | 1 | 1 | 1 | 0 | 1 | 1 | 2 | 0 | 2 | 0 |
| 2 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | 2 | 0 | 0 | 1 | 2 | 0 |
| 3 | 0 | 0 | 0 | 0 | 3 | 1 | 1 | 2 | 2 | 1 | 0 | 1 | 1 |
| 4 | 2 | 1 | 0 | 0 | 1 | 1 | 2 | 0 | 2 | 0 | 1 | 0 | 2 |

Showing 1 to 4 of 297 entries, 500 total columns

View RF vote for each case in the dataset via `m.RF.1$votes`

```
> m.RF.1$votes
```

| | No | Yes |
|----|------------|-------------|
| 1 | 0.56497175 | 0.435028249 |
| 2 | 0.10270270 | 0.897297297 |
| 3 | 0.06217617 | 0.937823834 |
| 4 | 0.55172414 | 0.448275862 |
| 5 | 0.98275862 | 0.017241379 |
| 6 | 0.96111111 | 0.038888889 |
| 7 | 0.28061224 | 0.719387755 |
| 8 | 0.61538462 | 0.384615385 |
| 9 | 0.16279070 | 0.837209302 |
| 10 | 0.30612245 | 0.693877551 |
| 11 | 0.65625000 | 0.343750000 |
| 12 | 0.80113636 | 0.198863636 |
| 13 | 0.32105263 | 0.678947368 |
| 14 | 0.75661376 | 0.243386243 |
| 15 | 0.82291667 | 0.177083333 |
| 16 | 0.84390244 | 0.156097561 |
| 17 | 0.72093023 | 0.279069767 |
| 18 | 0.82474227 | 0.175257732 |
| 19 | 0.90659341 | 0.093406593 |
| 20 | 0.98445596 | 0.015544041 |

Q: Consider case 1. Does this mean 56% of the 500 trees voted AHD = No and 44% of the 500 trees voted AHD = Yes?

Ans: No. Not 500 trees. Only in those trees (approx. 1/3 of 500) for which case 1 is OOB.

`votes` (classification only) a matrix with one row for each input data point and one column for each class, giving the fraction or number of (OOB) 'votes' from the random forest.

Caution: `m.RF.1$err.rate` is not the error rate at each tree

```
> err.rate <- m.RF.1$err.rate  
> view(err.rate)
```

| | OOB | No | Yes |
|---|-----------|-----------|-----------|
| 1 | 0.2654867 | 0.3015873 | 0.2200000 |
| 2 | 0.2696629 | 0.2755102 | 0.2625000 |
| 3 | 0.2590909 | 0.2333333 | 0.2900000 |
| 4 | 0.2701613 | 0.2388060 | 0.3070175 |
| 5 | 0.2613636 | 0.2291667 | 0.3000000 |
| 6 | 0.2637363 | 0.2482759 | 0.2812500 |
| 7 | 0.2588652 | 0.2432432 | 0.2761194 |
| 8 | 0.2465278 | 0.2287582 | 0.2666667 |

Q: Consider row 4. What is the meaning of OOB = 0.27? Does this mean the 4th tree OOB error is 27%?

Ans: OOB error using the first 4 trees is 27%

`err.rate`

(classification only) vector error rates of the prediction on the input data, the i-th element being the (OOB) error rate for all trees up to the i-th.

Python RF Implementation

- Categorical Y:
 - `sklearn.ensemble.RandomForestClassifier`
 - <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>
- Compared to R `randomForest()`:
 - Python used default $B = 100$. **Please overwrite to 500.**
 - Python also used default RSF size = $\text{int}(\sqrt{M})$.

Python RF Implementation

- Continuous Y:
 - `sklearn.ensemble.RandomForestRegressor`
 - <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestRegressor.html>
- Compared to R `randomForest()`:
 - Python used default $B = 100$. Please overwrite to 500.
 - Python used default RSF size = M . Please overwrite to $\text{int}(M/3)$.

Summary

- Random Forest
 - Bagging
 - 500 Bootstrap samples
 - Maximal CART per Bootstrap sample.
 - Random Subset Feature Selection
 - Default size $\text{floor}(M/3)$ for continuous Y
 - Default size $\text{floor}(\text{sqrt}(M))$ for categorical Y
 - Errors calculated from OOB cases
 - Check errors stabilised before reaching ntree (default 500).
 - Random Forest has a special way to assess variable importance.